

Your Proposal Title

Author 1 | SCIPER 1 | author1@epfl.ch

Author 2 | SCIPER 2 | author2@epfl.ch

Author 3 | SCIPER 3 | author3@epfl.ch

Author 4 | SCIPER 4 | author4@epfl.ch

Your group name

Abstract

A concise summary (under 300 words). State the gaps or motivation behind the work, your approach to producing the individual and group models, and the main empirical findings. Keep implementation details out of the abstract — those belong in the body.

1 Introduction

Motivate the problem. Sketch how current methods fall short (with references to prior work), give a high-level overview of your individual and group model strategies, and preview the headline results. Example citation: (Gusfield, 1997).

2 Method

Describe your training approach in enough detail that a reader could re-implement it. Cover all five models — the four individual models (one per benchmark) and the group model. For each, motivate the choice of post-training method (SFT, RLHF, RLVR, or a combination), formalize the training objective where it helps, and explain non-trivial implementation choices. Clarify how the group model is obtained from the individual models, distillation, model merging, or just how you balanced specialization against generalization. Use equations and figures wherever they aid clarity. You can create separate sections like the following:

2.1 Maths

2.2 Safety

2.3 Group Model

3 Experimental Setup

3.1 Datasets

For each model, describe the training data: sources, sizes, preprocessing steps, and any filtering or augmentation. Note licensing, legal, or ethical consid-

erations. Check if a summary table aids clarity for the reader.

3.2 Evaluation

Specify the metrics you report, the validation sets y(their construction, coverage, etc), and the baselines you compare against. Justify why each baseline is the right one to beat.

3.3 Implementation Details

Report hyperparameters, optimizer settings, training duration, hardware used, and anything else needed for reproducibility. Tables might help here again for clarity.

4 Results

4.1 Main Results

Present quantitative results with tables and plots. Compare each individual model against baselines on its dedicated benchmark, and report the group model's performance across all four benchmarks side-by-side.

4.2 Analysis

Go beyond reporting numbers. Discuss what worked and what did not like ablations of method components, comparisons across post-training methods, training instabilities, and unexpected failure modes. Null results are informative: when something did not work, say so, hypothesize why, and what you tried to tackle it. Close with the main takeaways the reader should remember.

5 Ethical Considerations

Discuss the broader impact of your work: who benefits, who could be harmed, and whether any harms fall disproportionately on marginalized groups. Suggest mitigations for the harms you identify (like adaptation to lower-resource languages), and argue why your work is, on balance, a positive contribution despite the possible harms.

6 Conclusion

Summarize your main contributions and findings, state the primary limitations of your work honestly, and suggest concrete directions for future investigation.

References

Dan Gusfield. 1997. *Algorithms on Strings, Trees and Sequences*. Cambridge University Press, Cambridge, UK.

A Appendix (optional)

This does not count towards your 4-page limit. You can put additional results, tables, and figures here; for example, extended hyperparameter sweeps, failed experiments, qualitative generation examples, or prompt templates. Your main body should still be self-sufficient, i.e., anything important to your argument should go there, and you should assume that you will be graded based on the content of the main part of your paper only.